Off the Hook

Disasters and the Absence of Accountability in Authoritarian Russia

Noah Buckley
Trinity College Dublin

September 1, 2026

When disaster strikes under autocracy

- Autocracies lean on **performance** to legitimate themselves
- Disasters are exogenous tests of that performance (Gerschewski 2013)
- When a disaster strikes, does the public hold the regime to account?
  - Democratic accountability: voters **punish** poor performance (Healy & Malhotra 2009; Gasper & Reeves 2011)
  - But Lazarev et al. (2014): the 2010 wildfires **raised** support for the authorities — a *rally*
- Both are claims about an **average effect**. Which one is right?

What we know and what we don't

- Rich work on authoritarian responsiveness: blame attribution (Reuter & Beazer 2019; Buckley et al. 2022), affective bonds (Greene & Robertson 2019; Frye 2021)
- A single, striking result on disasters: the Lazarev *rally* (Lazarev et al. 2014)
- But the record is **contradictory**, and almost all of it rests on:
  - one event, or one region, or
  - an *average* effect estimated without a clean identification check
- No multi-event, micro-level test of disaster accountability in an autocracy

Research question

1. Does exposure to a local disaster shift **support** for the regime (Putin) or its **agents** (governors)?

What I find

- The effect on **both** regime and governor approval is a **null**
- Apparent effects in naive models are **confounded**
- The design is **well powered** (MDE < 1pp) — so the null is informative, not a power failure
- **No moderator** survives: severity, salience, novelty, type, period, or respondent characteristics

Authoritarian publics do not hold the regime to account for disasters.

Data & design

Disaster data

- **4,235** disasters and accidents, 2012–2022, 84 regions
- Built from FPP regional digests (≈95,000 raw events)
  - hand-coding **+** LLM classification **+** rule-based and LLM filters
- Deliberately **broad**: fires, floods, industrial and mining accidents, transport, gas explosions, mass poisonings, infrastructure failures
- Each event: death toll, ordinal severity, ex-ante government responsibility, type
- **External validation**
  - recovers 96% of well-known Wikipedia-listed disasters
  - corroborates with the EM-DAT registry
  - 89% recall / 85% non-disaster catch on a hand-coded sample

Source: FPP regional digests, 2012–2022

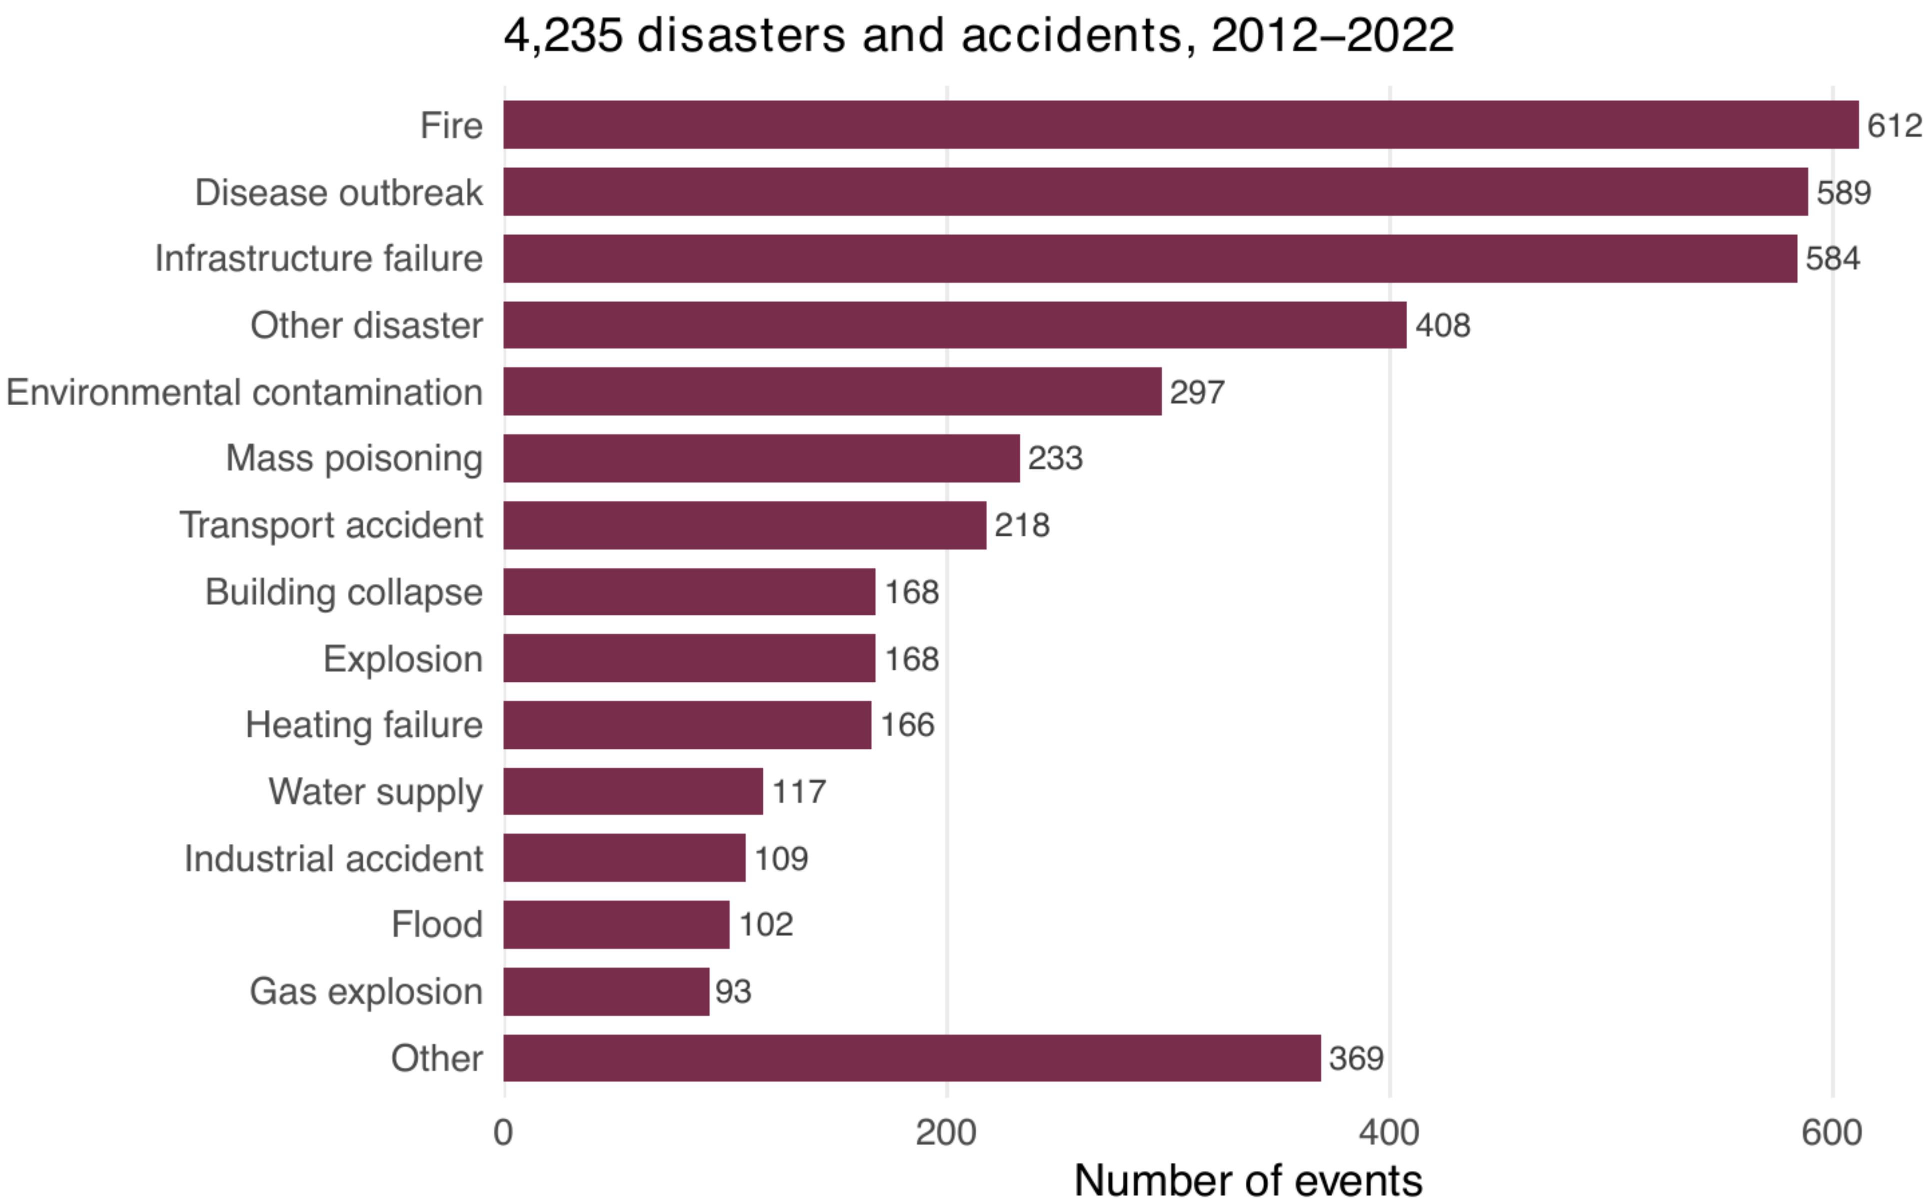
4,235 disasters and accidents, 2012–2022
Fire
612
Disease outbreak
589
Infrastructure failure
584
Other disaster
408
Environmental contamination
297
Mass poisoning
233
Transport accident
218
Building collapse
168
Explosion
168
Heating failure
166
Water supply
117
Industrial accident
109
Flood
102
Gas explosion
93
Other
369
0
200
400
600
Number of events

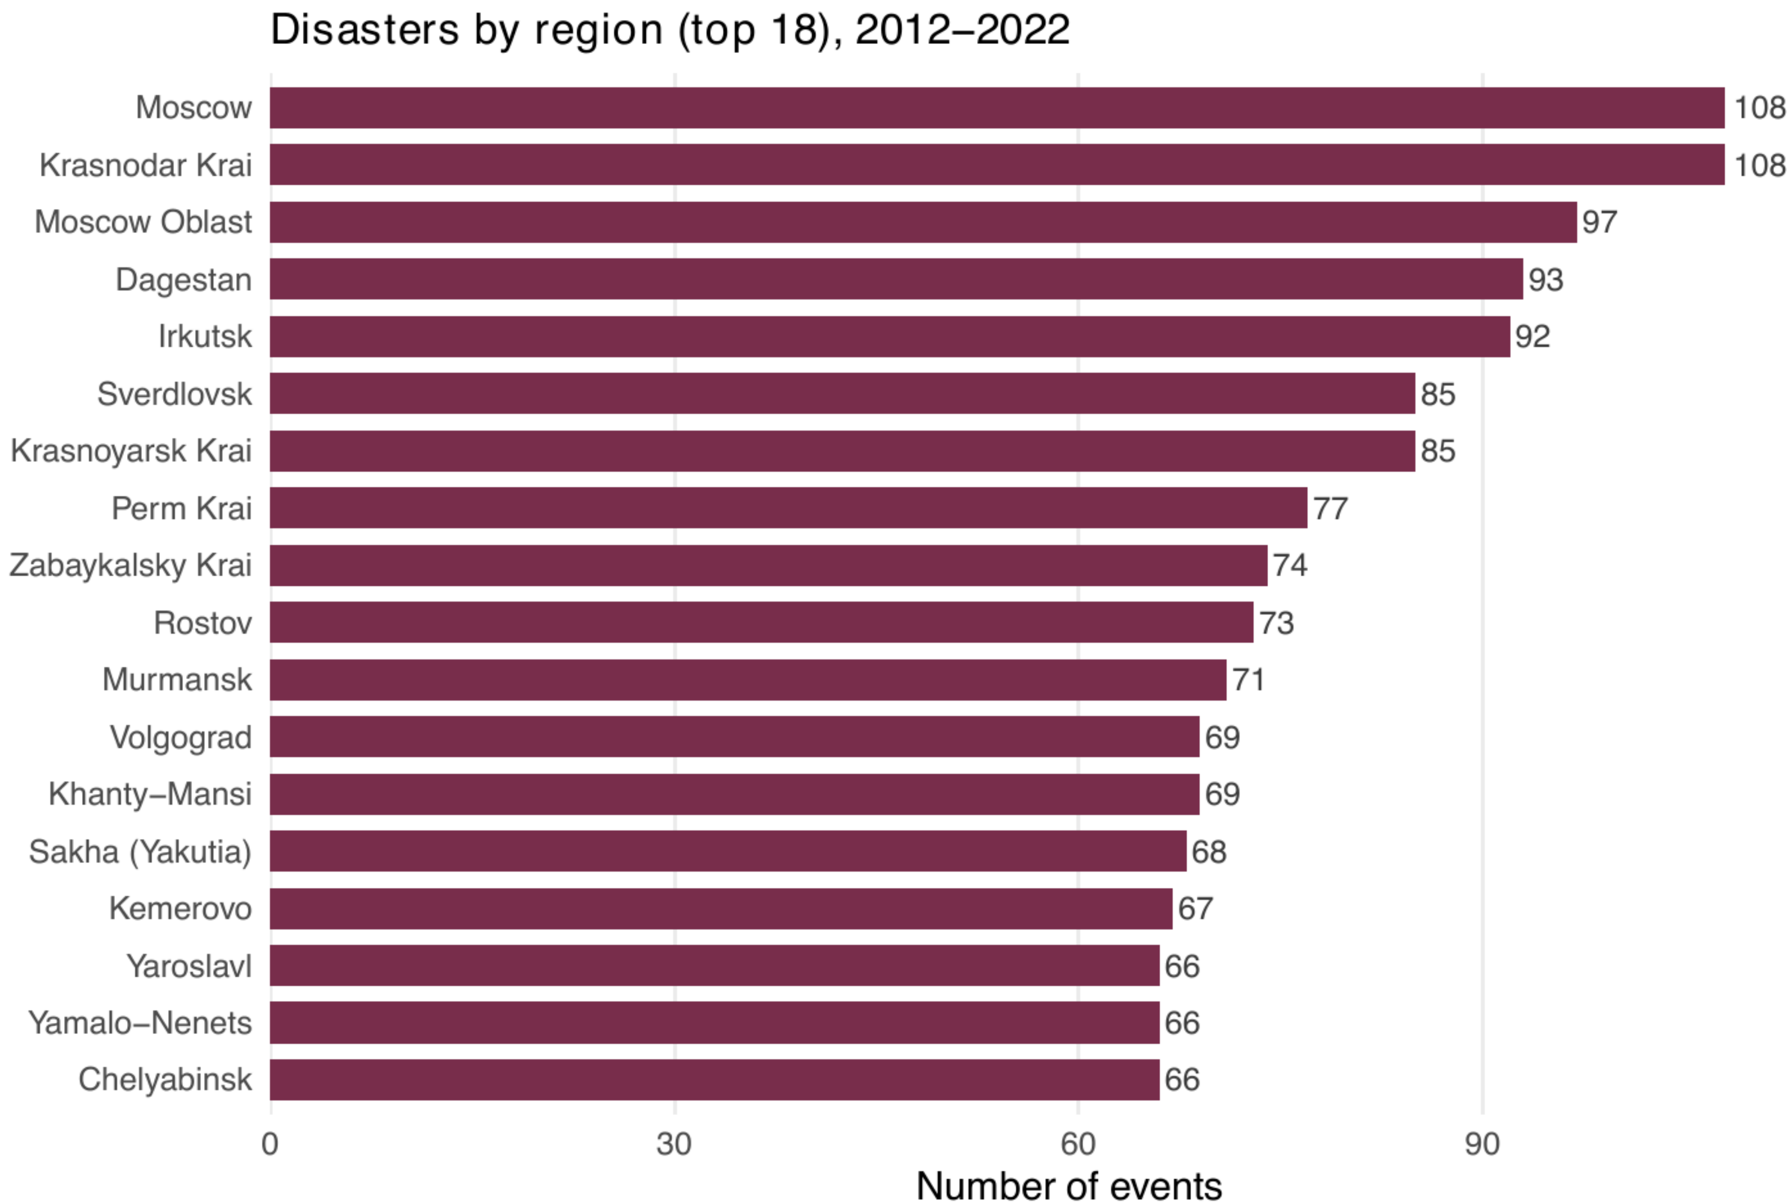
Disasters by region (top 18), 2012–2022
Moscow 108
Krasnodar Krai 108
Moscow Oblast 97
Dagestan 93
Irkutsk 92
Sverdlovsk 85
Krasnoyarsk Krai 85
Perm Krai 77
Zabaykalsky Krai 74
Rostov 73
Murmansk 71
Volgograd 69
Khanty–Mansi 69
Sakha (Yakutia) 68
Kemerovo 67
Yaroslavl 66
Yamalo–Nenets 66
Chelyabinsk 66
0
30
60
90
Number of events

Survey data

- Pooled, harmonized Russian public opinion data — Levada, VTsIOM, FOM
  - a harmonized database of 5.2 million responses, 1993–2024
  - restricting to the disaster window (2012+) and the needed covariates, the analysis uses **≈ 1.1 million** (Putin approval) and **≈ 330,000** (governor)
  - interview date and region for every respondent
- Outcomes (dichotomous):
  1. Approval of **Putin** — diffuse / regime-level support
  2. Approval of the **governor** — specific / lower-level support
- Plus the daily VTsIOM "Sputnik" survey (≈3.3M, 2018–2024) for high-frequency event studies

Estimation

Individual-level difference-in-differences

- Linear probability model — OLS with two-way fixed effects
- Fixed effects: **region, year-month**, survey source, locality size
- SEs clustered by region

Treatment: a disaster in the respondent's region in the last **14 / 30 days**

- Identification comes from **local exposure** — whether a disaster recently struck *your* region
- Parallel trends is untestable, so I rely on a **pre-event placebo**

The key diagnostic: a pre-event placebo

A disaster that has not happened yet cannot move current approval.

- Estimate the effect of a disaster in the **next** 30 days — a *forward* placebo
- If the “forthcoming” disaster predicts approval as well as the past one, the estimate reflects **selection on which region-months get recorded** as having disasters

Results

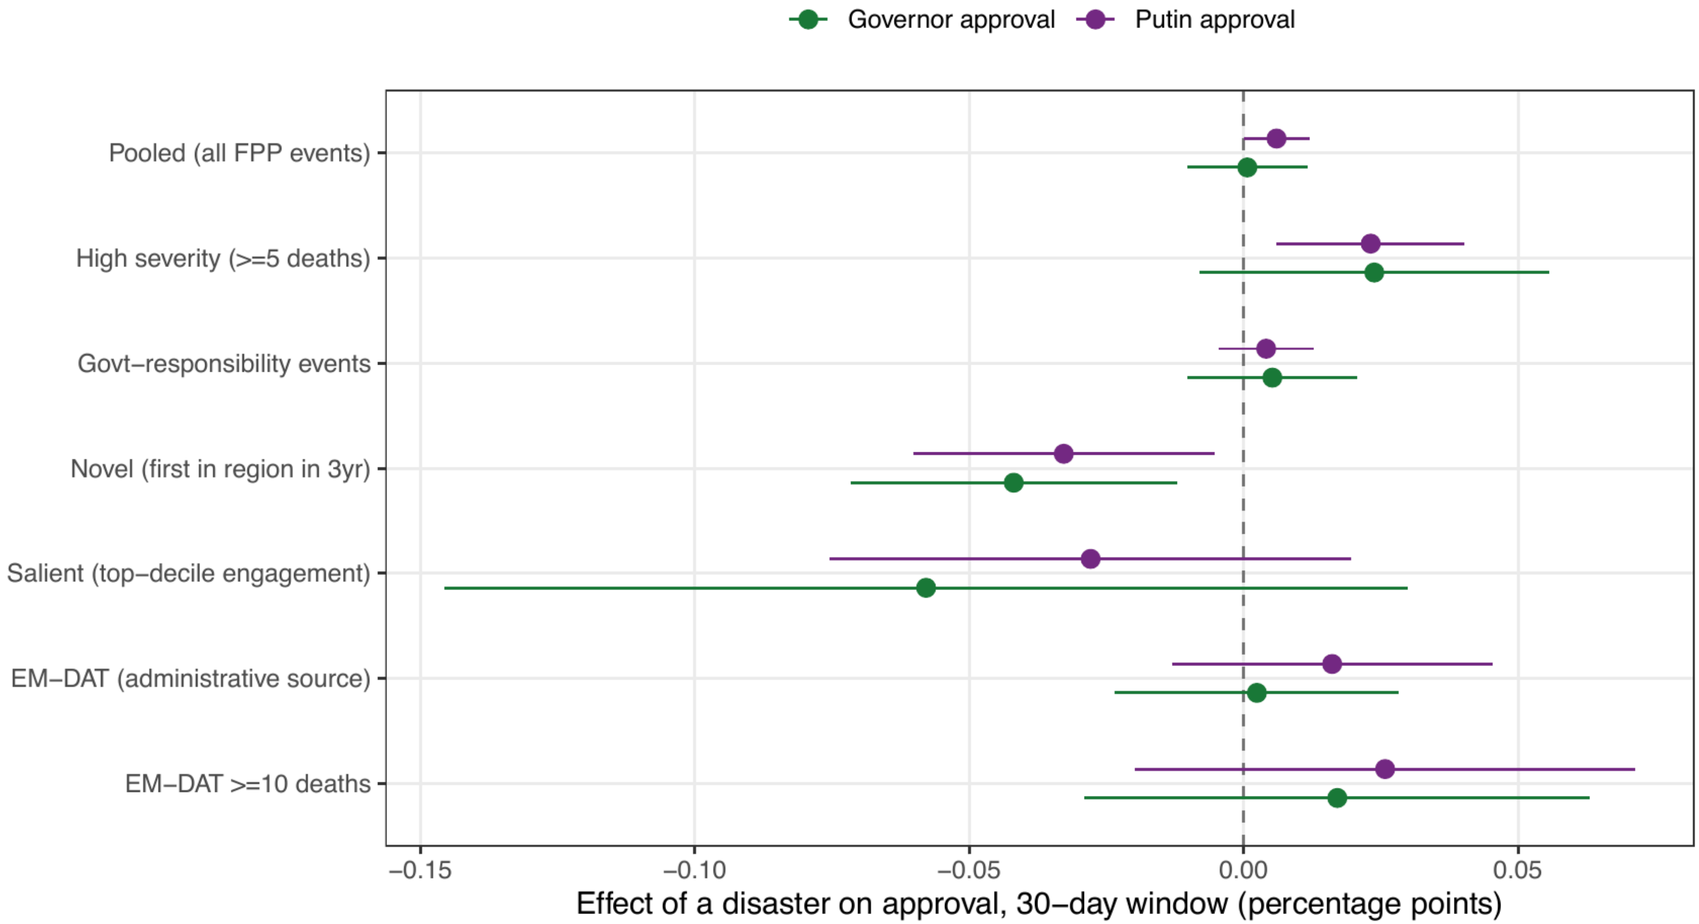
Governor approval
Putin approval
Pooled (all FPP events)
High severity (>=5 deaths)
Govt−responsibility events
Novel (first in region in 3yr)
Salient (top−decile engagement)
EM−DAT (administrative source)
EM−DAT >=10 deaths
−0.15
−0.10
−0.05
0.00
0.05
Effect of a disaster on approval, 30−day window (percentage points)

Results: Null effects

DV: approval. "Placebo" = a disaster in the next 30 days. SEs clustered by region.

	Putin: effect	Putin: placebo	Govr: effect	Govr: placebo
Disaster, 30d	**0.006**$^{+}$	0.004	0.001	0.005
	(0.003)	(0.004)	(0.006)	(0.006)
Region / month / source FE	✓	✓	✓	✓
N	1,066,957	1,066,957	329,988	329,988

- The one marginally-positive coefficient (Putin, +0.6pp) is **matched by its pre-event placebo** (+0.4pp)

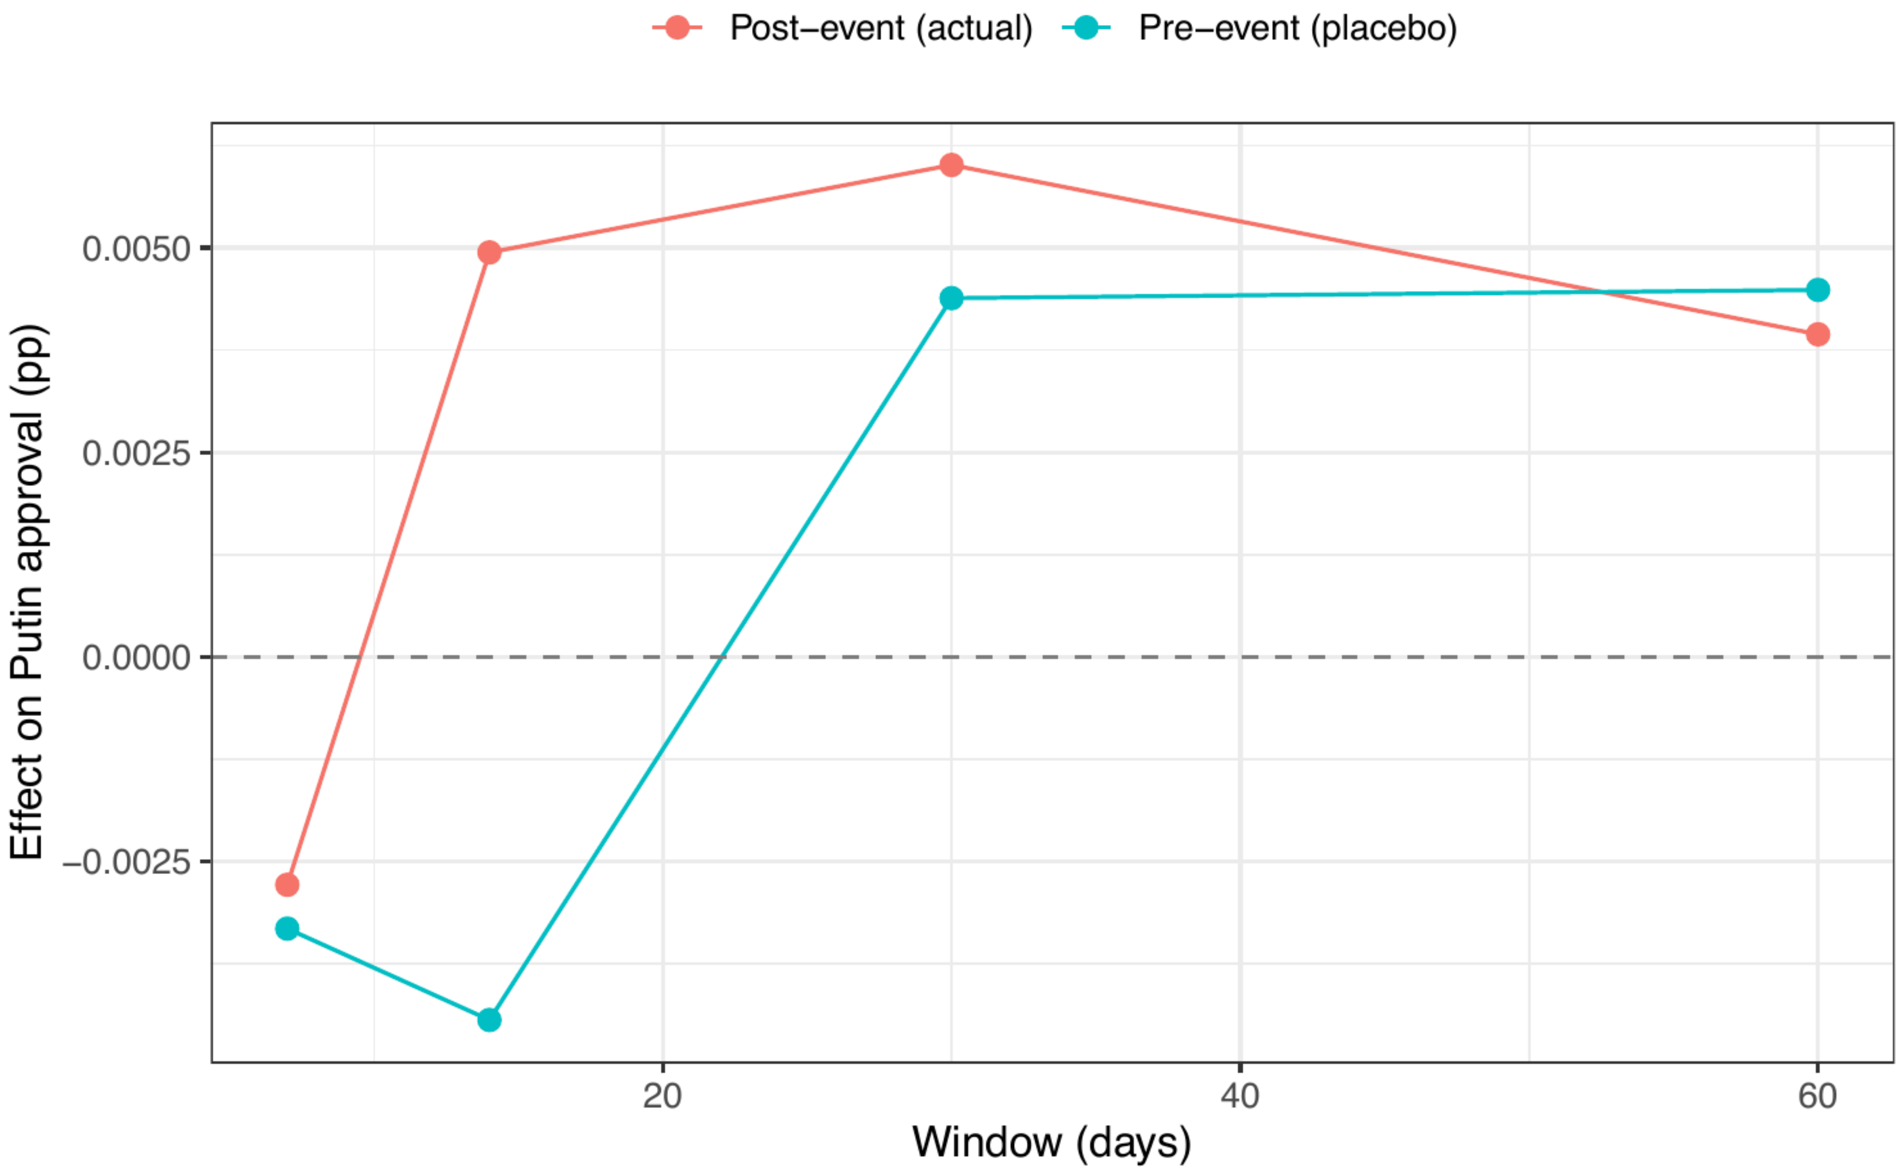
Post−event (actual)
Pre−event (placebo)
Effect on Putin approval (pp)
0.0050
0.0025
0.0000
−0.0025
20
40
60
Window (days)

Naive estimates are confounded

- Apparent positive effects appear **only with wide windows** (14–30 days)
- They **vanish** under region × year fixed effects
- Pre-event placebo is the same size

The “rally” and the “punishment” are the same confounded association, sliced two ways.

- Plausibly the attention of media compilers is already trained on these region-months — the *disasters* are exogenous, but their **recording** is not

Cross-source confirmation

- The FPP source is media-selected — so I re-estimate with **administrative** sources:
  - **EM-DAT** (international registry; deaths/affected/emergency thresholds)
  - **EMERCOM** (emergency-service counts)

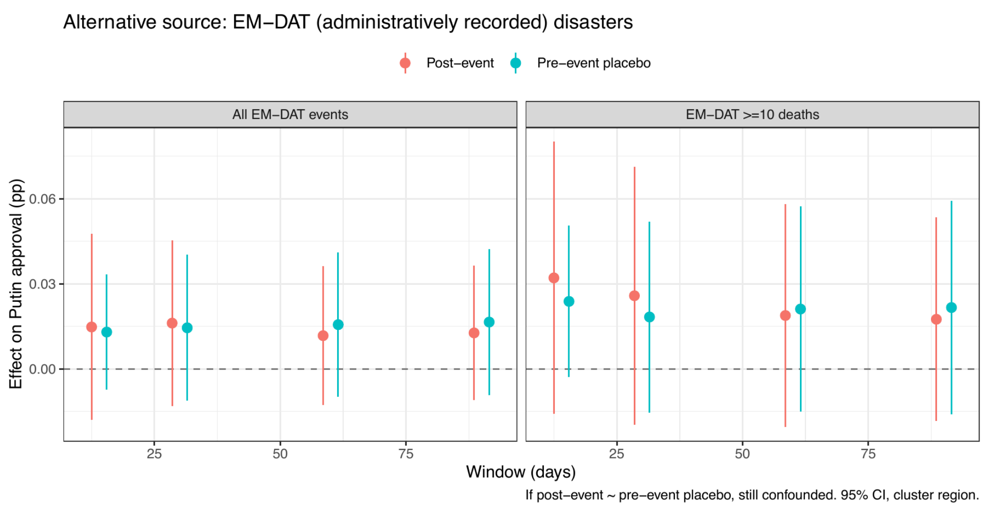

- Both are far less attention-driven. Both are **null**, with the pre-event placebo tracking the estimate.

No moderator survives

The pre-registered moderators **and** an extended battery:

- **Severity** thresholds (deaths ≥ 5, ≥ 10) — confounded, not causal
- **Salience** — eight measures of local vs. national attention
- **Novelty** — six ways of coding "surprise"
- **Event type** — fires, floods, industrial, transport, … one at a time
- **Period** — performance vs. ideological legitimation eras
- **Respondents** — age, urban/rural, internet vs. TV

Nothing produces a robust effect that survives the placebo.

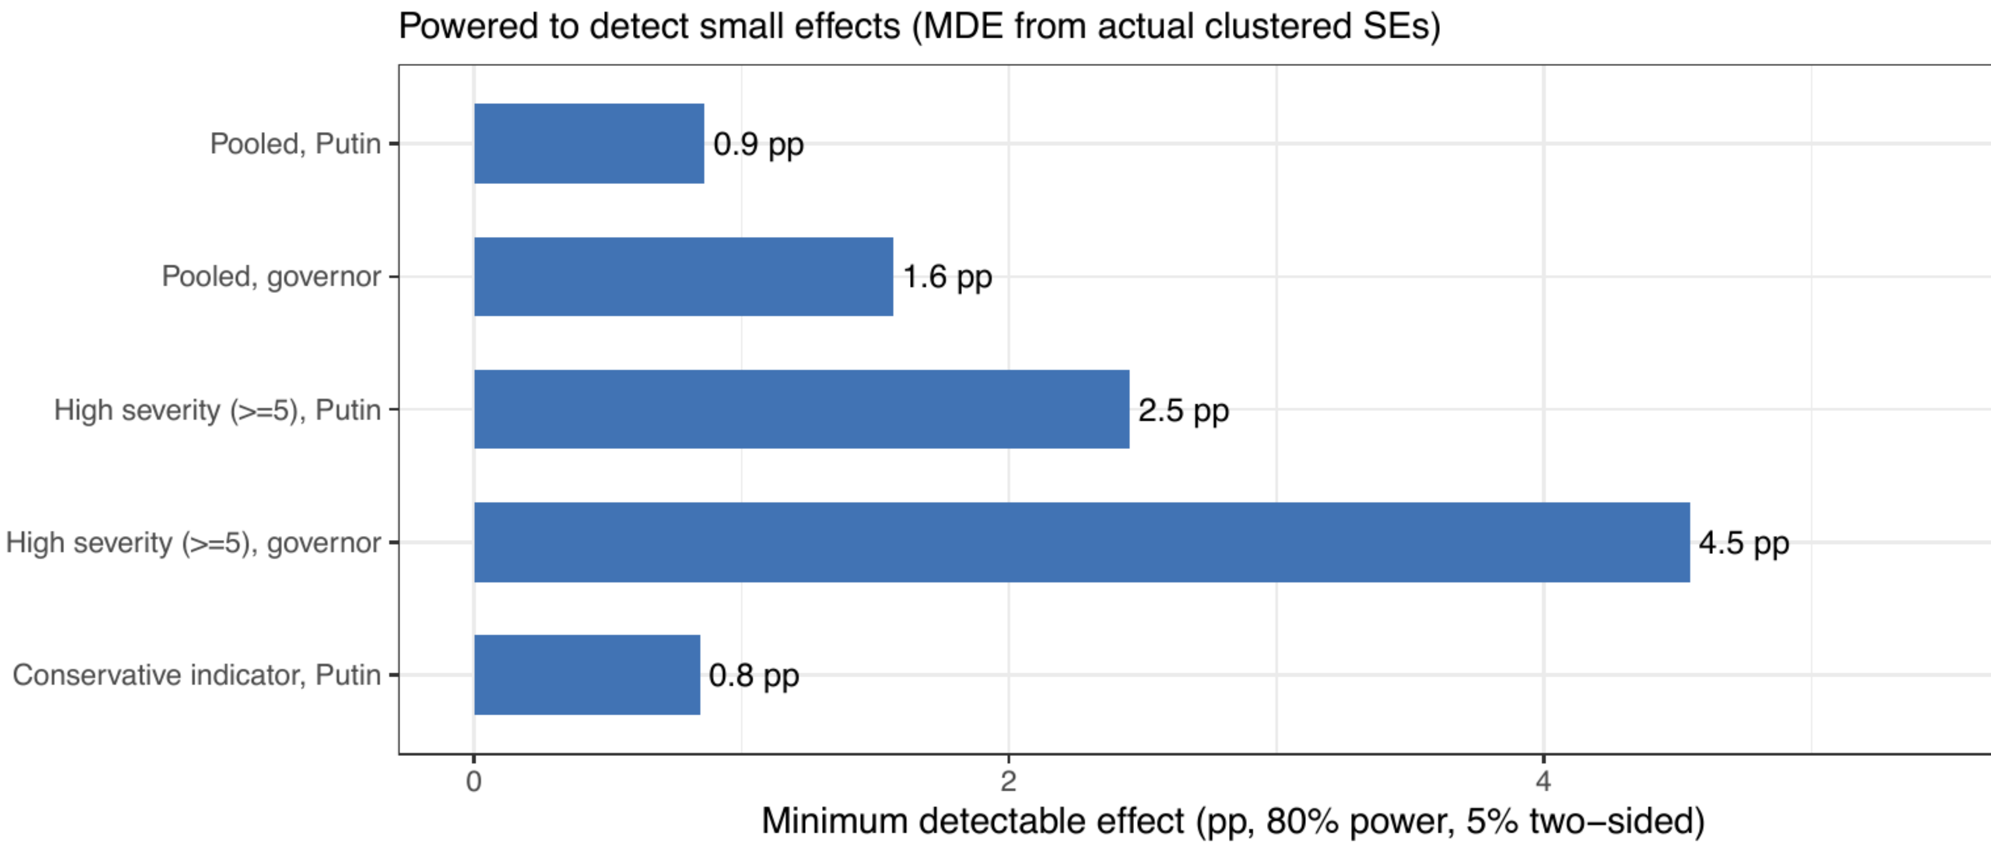
Powered to detect small effects (MDE from actual clustered SEs)
Pooled, Putin
Pooled, governor
High severity (>=5), Putin
High severity (>=5), governor
Conservative indicator, Putin
0.9 pp
1.6 pp
2.5 pp
4.5 pp
0.8 pp
0
2
4
Minimum detectable effect (pp, 80% power, 5% two–sided)

Are we just underpowered? No

- Minimum detectable effect from the **actual** clustered SEs:
  - **< 1pp** on the pooled Putin specifications
  - ≈ 2.5pp on the severity-restricted cuts
- We would comfortably detect modest effects — and we find nothing

The null is informative, not a power failure.

Legitimation and Lazarev

- The event data begin in 2012; surveys and EM-DAT reach back further
- In the **earlier, performance-legitimation period**, severe events show a **rally** — directionally consistent with Lazarev et al. (2014) (Lazarev et al. 2014)
- But this does **not generalize**:
  - timing is episodic, not a clean break
  - does not survive comprehensive, FPP-comparable event measurement

Why is there no accountability?

- A plausible reading: the **blame channel is severed**
  - heavy state control of the information environment
  - a regime that legitimates increasingly through ideology, not performance
- The public clearly **notices** disasters; approval simply does not move — attention without accountability
- I read this *tentatively*. It motivates mechanism work; it is not tested here.

Conclusions

- No measurable accountability for disasters in Putin-era Russia — even severe, government-implicated ones — across **three** data sources
- Naive estimates mislead: a “rally” and a “punishment” are confounded artifacts of event recording
- A methodological caution for observational disaster–opinion designs: **run the pre-event placebo**

Future work

- Direct mechanism tests (internet vs. TV; media framing)
- A continuous measure of regime framing (performance vs. ideology)
- High-salience post-2024 cases (e.g., the Kerch oil spill)

Thank you

Noah Buckley

noah.buckley@tcd.ie · www.noahbuckley.me

Appendix

High-severity events: effect vs. placebo

Deaths ≥ 5. The larger Putin estimate is matched by a comparable pre-event placebo.

	Putin: effect	Putin: placebo	Gov: effect	Gov: placebo
High-severity, 30d	**0.023****	0.014^{+}	0.024	0.000
	(0.009)	(0.008)	(0.016)	(0.016)
Region / month / source FE	✓	✓	✓	✓
N	1,066,957	1,066,957	329,988	329,988

$^{+}$ $p<0.1$, ** $p<0.01$

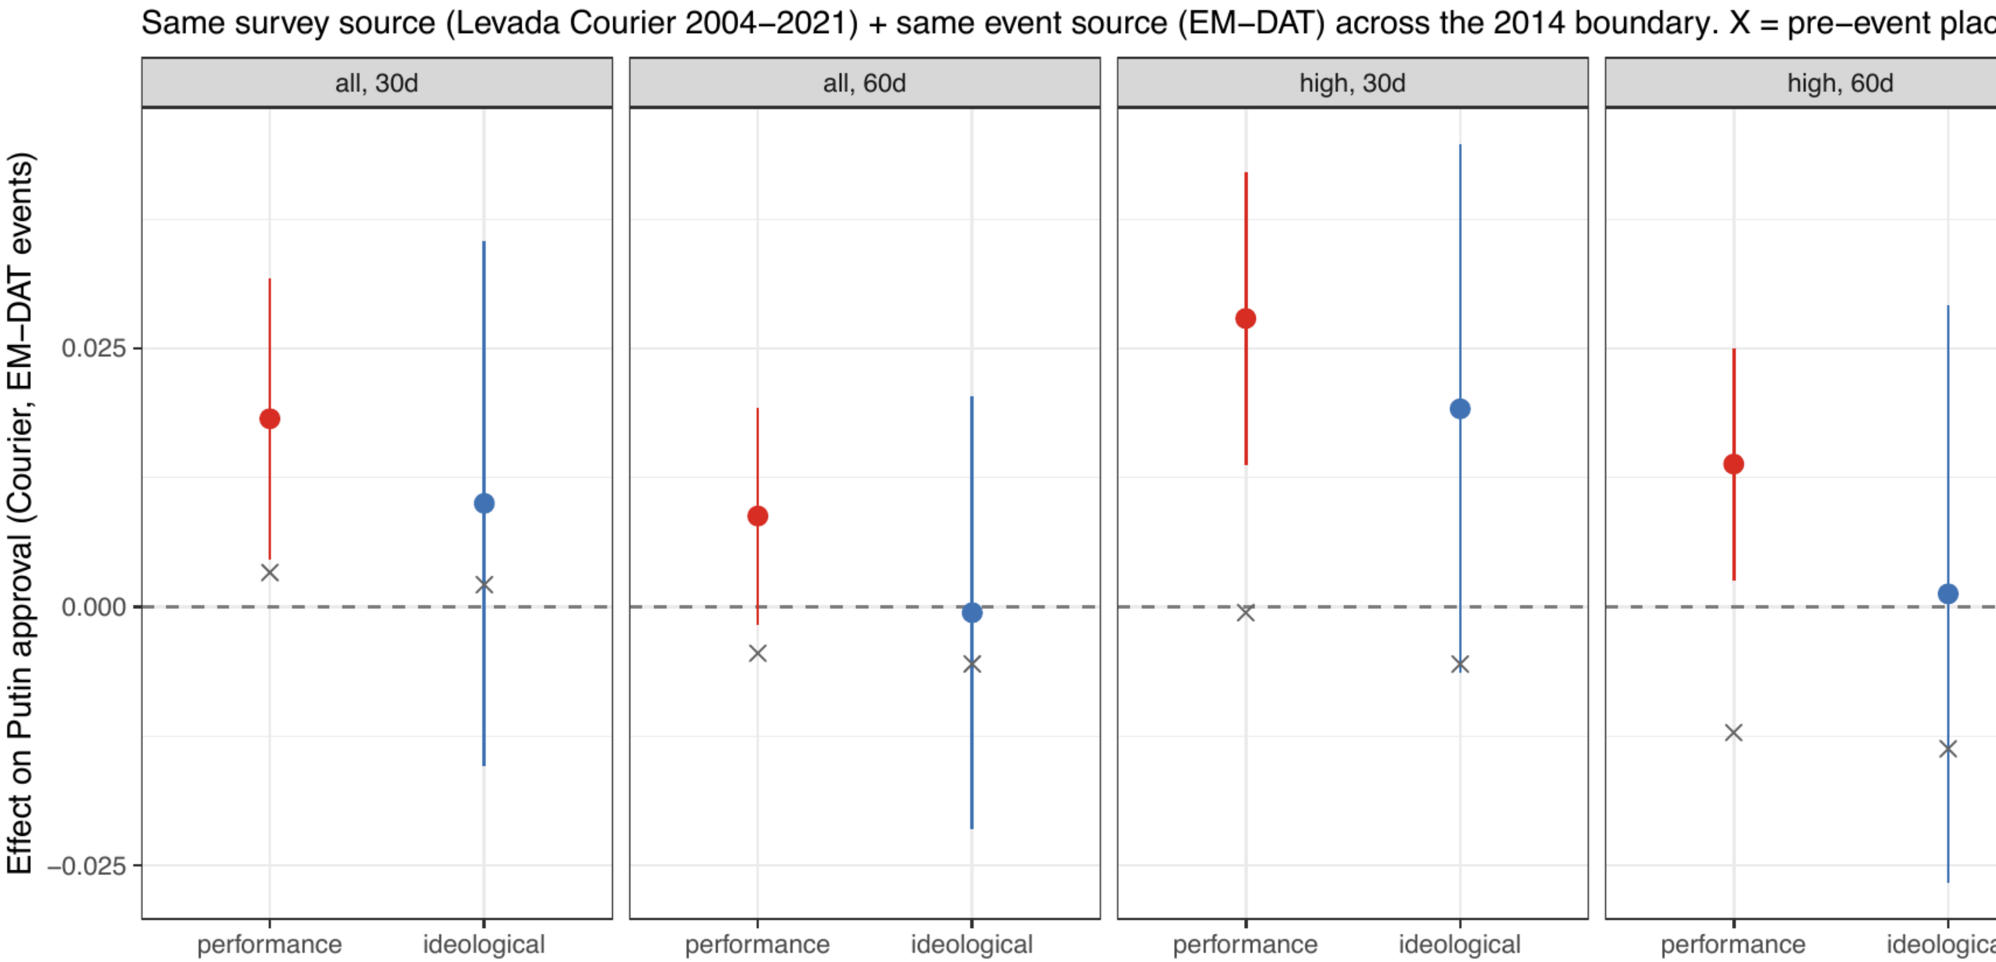
Same survey source (Levada Courier 2004–2021) + same event source (EM–DAT) across the 2014 boundary. X = pre–event plac
all, 30d
all, 60d
high, 30d
high, 60d
Effect on Putin approval (Courier, EM–DAT events)
0.025
0.000
–0.025
performance
ideological

Cluster-robust inference

- 84 region clusters; CR1 inference is adequate
- Two-way (region × month) clustering does **not** weaken the marginal Putin positive — it *sharpens* it ($p \approx 0.05 \to 0.01$)
- So the null causal reading rests on the **placebo**, not the standard error: a precisely-estimated positive matched by an equally-positive pre-event placebo is *more* clearly selection